



UNIVERSITÀ DI PISA

**Computer Engineering
Electronics and Communications Systems**

Finite Impulse Response filter
with SRRC impulsive response

Project Documentation

Matteo Guidotti

Academic Year 2020/2021

Contents

1	Introduction	3
1.1	Specifications	3
1.2	The intersymbol interference	3
1.3	The SRRC filter	4
1.4	Finite Impulse Response filters	4
2	Architecture Description	5
3	Test Plan	6
3.1	Coefficients Quantization	7
3.2	Binary Conversion	7
3.3	Signal Generator	7
3.4	FIR Software Simulation	8
3.5	VHDL Testbench	8
3.6	Result Checker	9
3.7	Results Comparison	9
4	Synthesis and Implementation	10
4.1	Warning Analysis	11
4.2	Timing Summary	11
4.3	Area Occupation	12
4.4	Power Summary	14
4.5	Utilization Summary	14
5	Conclusions	14

1 Introduction

1.1 Specifications

This project is about the implementation of a **Finite Impulse Response filter (FIR)** characterized by a **Square Root Raised Cosine (SRRC)** impulsive response. The filter is characterized by:

- Filter order equal to $N = 22$
- Number of samples per symbol = 4
- Roll-Off factor = 0.5
- The coefficients are:

$$\begin{aligned} c_0 = c_{22} &= -0.0165; c_1 = c_{21} = -0.0150; c_2 = c_{20} = 0.0155; c_3 = c_{19} = 0.0424; \\ c_4 = c_{18} &= 0.0155; c_5 = c_{17} = -0.0750; c_6 = c_{16} = -0.1568; c_7 = c_{15} = -0.1061; \\ c_8 = c_{14} &= 0.1568; c_9 = c_{13} = 0.5786; c_{10} = c_{12} = 0.9745; c_{11} = 1.1366. \end{aligned}$$

$$y[n] = \sum_{i=0}^N c_i \cdot x[n - i]$$




Figure 1: Filter scheme

1.2 The intersymbol interference

In telecommunication, **intersymbol interference (ISI)** is a form of distortion of a signal in which one symbol interferes with subsequent symbols. This is an unwanted phenomenon as the previous symbols have a similar effect as noise,

thus making the communication less reliable. The spreading of the pulse beyond its allotted time interval causes it to interfere with neighboring pulses. ISI is usually caused by multipath propagation or the inherent linear or non-linear frequency response of a communication channel causing successive symbols to blur together. The presence of ISI in the system introduces errors in the decision device at the receiver output. Therefore, in the design of the transmitting and receiving filters, the objective is to minimize the effects of ISI, and thereby deliver the digital data to its destination with the smallest error rate possible.

1.3 The SRRC filter

The **Square Root Raised Cosine filter (SRRC)**, sometimes known as **Root Raised Cosine filter**, is frequently used as the transmit and receive filter in a digital communication system to perform matched filtering. This helps in minimizing intersymbol interference (ISI). The combined response of two such filters is that of the **raised-cosine filter**, in fact it obtains its name from the fact that its frequency response is the square root of the frequency response of the raised-cosine filter. To have minimum ISI, the overall response of transmit filter, channel response and receive filter has to satisfy **Nyquist ISI criterion**. The raised-cosine filter is the most popular filter response satisfying this criterion.

Summing up, the inter-symbol interference (ISI) deteriorates the received symbol quality in a band-limited digital communication system, but as I mentioned before it is well known that when a pair of matched SRRC filters are employed in the transmitter and the receiver, the overall frequency response of the system has a raised-cosine spectrum such that the ISI is gone at the time when the symbol is sampled. Unfortunately, ideal SRRC filters are not realizable. In reality, they can only be approximately implemented. In particular, such filters can be approximated with **Finite Impulse Response (FIR)** filters.

1.4 Finite Impulse Response filters

A **Finite Impulse Response (FIR) filter** is a filter whose impulse response (or response to any finite length input) is of finite duration, because it settles to zero in finite time. This is in contrast to **Infinite Impulse Response (IIR) filters**, which may have internal feedback and may continue to respond indefinitely (usually decaying). FIR filters can be discrete or continuous time, in our case we'll use the first one. For such a filter of order N , each value of the output is a weighted sum of the $N + 1$ most recent input values, where the weights are the $N + 1$ coefficients of the filter. FIR filters have also other important characteristics:

- They require no feedback. This means that any rounding errors are not compounded by summed iterations. The same relative error occurs in each calculation.

- They are inherently stable, since the output is a sum of a finite number of finite multiples of the input values.
- They can easily be designed to be linear phase by making the coefficient sequence symmetric. This property is sometimes desired for phase-sensitive applications, for example data communications, seismology, crossover filters, and mastering.
- The main disadvantage of FIR filters is that considerably more computation power in a general purpose processor is required compared to an IIR filter with similar sharpness or selectivity.

2 Architecture Description

The FIR filter has been implemented by means of the VHDL code language and it is available in the file *VHDL/src/srrc_fir.vhd*.

The filter is characterized by three input and an output signals, that are the clock and the reset signals (1 bit) and x , that is the signal to be filtered (16 bits), besides y , that is the result of the filtering (16 bits).

```
-- Entity description
entity srrc_fir is
  port(
    clk : in std_logic; -- clock signal
    rst : in std_logic; -- reset signal
    x   : in std_logic_vector(15 downto 0); -- Input signal
    y   : out std_logic_vector(15 downto 0) -- Output signal
  );
end srrc_fir;
```

Figure 2: Entity structure

The computation of the result requires to use the last 23 signal samples, so it is necessary to store 22 samples, because the most recent one is simply taken by the input signal. In order to store them, 22 D-flip-flops edge triggered are instantiated and at every clock rising edge they will be uploaded with the correct value, with the respect to their positioning. The DFFs are implemented using a synchronous active low reset. Given that the multiplier circuits are very expensive from the power and time efficiency point of views, I decided to exploit the symmetry of the coefficients to instantiate only 12 multipliers, instead of 23. So, before the multiplier circuits, every couple of samples that must be multiplied with the same coefficient are added up together and each output will be then multiplied by the related coefficient. This is true for all the stored samples, except for the one that must be multiplied with the c_{11} coefficient, that must not be summed to any other sample. Then, all the outputs of the multiplications must be summed up together

in order to obtain the final result. After the simulation of the testbench I decided to represent the results of the first sums with 17 bits in order to be sure to not loose any information, and I applied the same reasoning to the results of the multiplications, so that the related outputs are represented using 33 bits. I decided, instead, to not consider the carry out of the last sum because I noticed that any information would have been lost. At the end, the final output must be represented with 16 bits, because of the specification; by observing the result of the testbench I decided that the best configuration is the one in which the 16 LSBs of the output of the final sum are truncated and the MSB is not taken into account, so that a saturation of 1 bit is achieved. In order not to have stability problems, the output is synchronized thanks to a DFF.

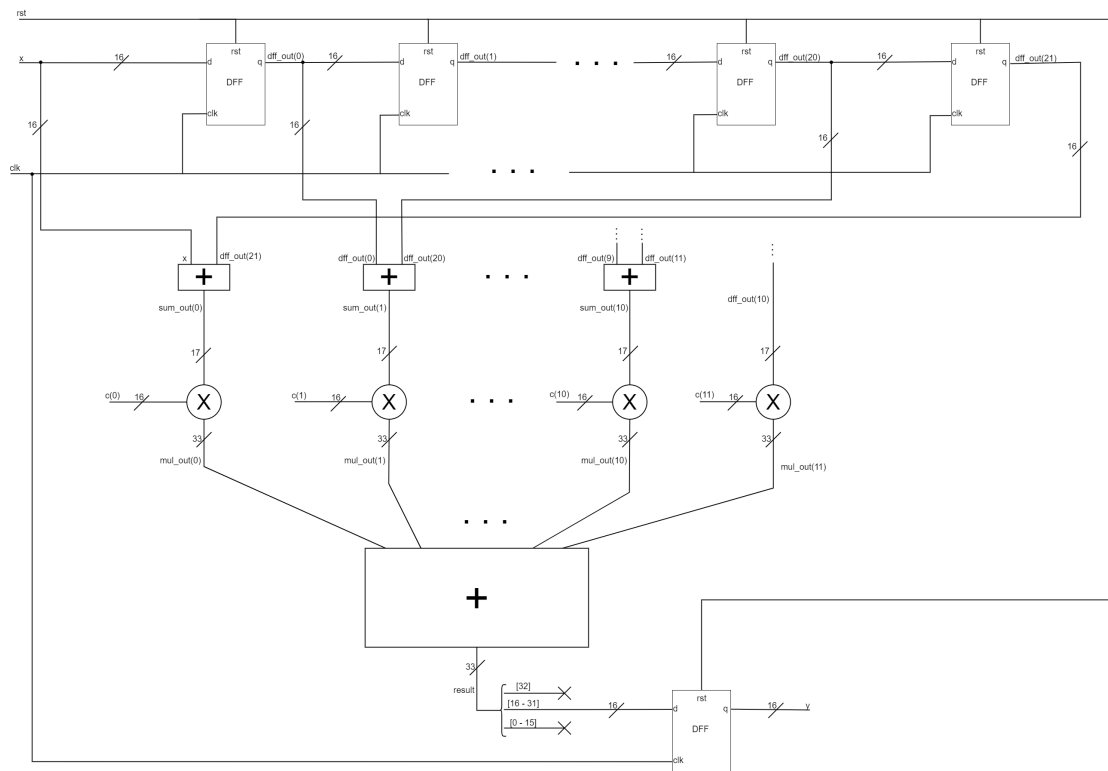


Figure 3: Architecture schema

3 Test Plan

In order to test the correct functioning of the implemented filter, I decided to simulate a sin waveform as the input signal. Now I will present all the utilities that I implemented so that I could be able to perform such a test.

3.1 Coefficients Quantization

The coefficients must be represented with 16 bits and they are all real numbers, so it is necessary to quantize them with a LSB in order to be represented in a fixed-point manner. So the values that will be represented on 16 bits are obtained by dividing each coefficient by the LSB and applying the floor function to the result. As a result, the quantization function is:

$$Q[x] = \text{floor}\left(\frac{x}{LSB}\right) \cdot LSB$$

For the sake of simplicity, I decided to use a LSB that is a power of 2, and given that all the coefficients are contained in the set $[-0.1568; 1.1366]$ I compute the mean quantization errors by using LSBs that were as small as possible and, at the same time, that return a result representable with 16 bits. The result of this study is that the smallest quantization error is obtained when the LSB is 2^{-14} , which consent to represent them with a mean error equal to $3.38 \cdot 10^{-5}$. The functions contained in the file *coefficient_quantization.py* compute the quantized values and print them in two different files. The file *coefficients.txt* will contain the values as they are, one for each row; while *coefToCode.txt* will contain the result of the computations in a format so that it consents me to copy and paste it directly to the vhd code. Then, the python utility compute also the mean quantization error for the just computed values.

3.2 Binary Conversion

The *binary_conversion.py* file contains a function that, given as an input an integer number, returns as a result its binary two's complement representation on 16 bits.

3.3 Signal Generator

The *signal_generator.py* file is responsible for generating an input for the filter. It generates a sin waveform signal computed on a customizable number of frequencies contained between 0 and the period of the waveform. The utility contains also a function that, given an LSB, compute the quantization for the signal values and the related quantization error. I decided to use a power of 2 LSB in this case too and, using the formula below, I found that the best LSB to represent values on $N = 16$ bits in the set $[-1; 1]$ is 2^{-15} , obtaining a mean quantization error equal to $1.53 \cdot 10^{-5}$.

$$LSB = \frac{|\min x|}{2^{N-1}}$$

The quantized values are printed in two different files in two different formats.

The file *quantized_signal.txt* will contain the quantized values as they are, one for each row; while *signal_to_code.txt* will contain the quantized values that are at first converted to a 16 bits representation and then encoded in a way so that I'm able to copy and paste them directly to the testbench file in the case statement that consents to the input signal to vary.




Figure 4: Example of 100 values between 0 and 2π

3.4 FIR Software Simulation

In order to check the correctness of the output values returned by the implemented filter in vhdl, I decided to implement the *fir_simulation.py* utility, that simulates in software the computations that the filter must perform. This script will request the input signal to the same utility that generates it for the testbench and then it will calculate the sequence of output values. The results are printed to the file *fir_sw_simulation_output.txt*: in each row of this file, a concatenation of the last input and the result of the previous computation is printed. This format is due to the fact that, in the hardware filter, the output is delayed with the respect of the input and so I simulated this behaviour also when the software implementation has to print them, such that the comparison between the results is simpler.

3.5 VHDL Testbench

The file *srrc_fir_tb.vhd* is the testbench file used in ModelSim to execute my filter implementation and check its behaviour. At first all the flip-flops are reset, then

the reset signal is driven high in order to consent the regular functioning of the DFFs. Starting from the third clock cycle, the input changes its values every clock cycle until the 103th cycle, in which the simulation is stopped. The case statement that regulates this behaviour is returned by the signal generator utility, as I already mentioned. At the end of the simulation, it is necessary to export an event list file that reports all the output values returned by the filter in order to check them respect to the software simulation results.

3.6 Result Checker

Once we obtained the results from both the software and the hardware simulations of the filter it is possible to check the correctness of the values obtained through the testbench simulation. The *result_checker.py* utility is responsible for taking the two files containing the outputs of the respective simulations and compare them. In order to perform a fair comparison, the script requires the LSB that must be used to interpret the values outputted by the testbench execution. The execution of the result checker will write one file and print to the standard output the mean error value obtained by the comparison of each result returned by the two simulations. Meanwhile, the file *results_comparison.txt* will contain, for each row, the concatenation of the software output and the relative hardware one, in order to have an immediate comparison between the two executions.

```
HW = 0.98388671875 <-> SW = 0.9842837167845508
HW = 1.23779296875 <-> SW = 1.2382172138292906
HW = 1.476806640625 <-> SW = 1.4772199765728407
HW = 1.705078125 <-> SW = 1.705517701872076
HW = 1.927490234375 <-> SW = 1.9279309800923936
HW = 2.144775390625 <-> SW = 2.1452703349005446
HW = 2.3543701171875 <-> SW = 2.354954513984524
```

Figure 5: Snippet of *results_comparison.txt*

3.7 Results Comparison

Thanks to these python scripts I have been able to check that the behaviour of my implementation of the FIR filter is able to return the correct values. As I already mentioned, the final result is represented with 33 bits, but the specifications consent me to use only 16 bits for the output signal. Given that the coefficients were

represented using a LSB of 2^{-14} and the input signal exploited a LSB equal to 2^{-15} , the resulting LSB for the final result is 2^{-29} , due to the multiplications between input signal samples and coefficients. This fact means that I have a maximum of 4 bits to represent the integer part of the results, but by watching at the results I observed that they were always contained in $(-4; 4)$, so it made me think that I could try to saturate the result by discarding the most significant bit.

I performed the testbench simulation at first by obtaining the outputs discarding the 17 least significant bits of the final results and I measured a mean error value of $4.67996 \cdot 10^{-4}$. On the other hand, the simulation has been also performed by truncating 16 least significant bits and discarding the most significant one, obtaining a mean error value of $4.493246 \cdot 10^{-4}$. Before executing *result_checker.py* is important to tune in a correct way the LSB used to interpret the output values of the hardware simulation, that depends on how many bits are truncated. When 17 bits are truncated, it will be equal to 2^{-12} , while in the other case it will be 2^{-13} . Said that, it is clear that both the two results are very good from the point of view of the precision of the results and it would occur some additional information about the input signals to define which bits are better to discard. In fact, if the signal is often characterized by values near the 0, it would be better the saturation scenario, while if the signal is often near, in absolute value, to 1, it would be better to choose the truncation only method. However, given the very small but observable difference between the two mean error values, I decided to use the configuration in which the most significant and 16 least significant bits are discarded.

4 Synthesis and Implementation

Before performing the synthesis, the implementation of a vhdl wrapper is necessary and it has been done in the file *srrc_fir_wrapper.vhd*. Vivado tool presents us the filter in this way:




Figure 6: Schematic of the filter

4.1 Warning Analysis

The only warning that popped up has arisen during the synthesis phase, that is:



Figure 7: Synthesis warning

Let's start by understanding what is the *Parallel synthesis*: from the <https://support.xilinx.com/> webpage we can read:

"In order to improve overall runtime of the design, Vivado synthesis launches parallel processes if the design is large enough and can benefit from the parallel flow. The parallel flow partitions the design into smaller 'RTL Partitions' which are processed independently by parallel processes. Vivado Synthesis will decide to use parallel flow only if the design size is large enough."

So the warning simply communicates that the design is not large enough to benefit from parallel synthesis. As a result, the parallel flow will not be run. This fact will not be a problem for the characteristics and the performance of my implementation.

4.2 Timing Summary

By imposing the constraint in which the clock has a period of $25ns$, the timing summary at the implementation stage is the following:

Design Timing Summary		
Setup	Hold	Pulse Width
Worst Negative Slack (WNS): 2,081 ns	Worst Hold Slack (WHS): 0,092 ns	Worst Pulse Width Slack (WPWS): 11,520 ns
Total Negative Slack (TNS): 0,000 ns	Total Hold Slack (THS): 0,000 ns	Total Pulse Width Negative Slack (TPWS): 0,000 ns
Number of Failing Endpoints: 0	Number of Failing Endpoints: 0	Number of Failing Endpoints: 0
Total Number of Endpoints: 788	Total Number of Endpoints: 788	Total Number of Endpoints: 238
All user specified timing constraints are met.		

Figure 8: Timing summary

As a consequence, the theoretical maximum usable frequency is:

$$f_{max} = \frac{1}{(25 - 2.081) \cdot 10^{-9}} \approx 43.6 \text{ MHz}$$


The resulting critical path is the following:

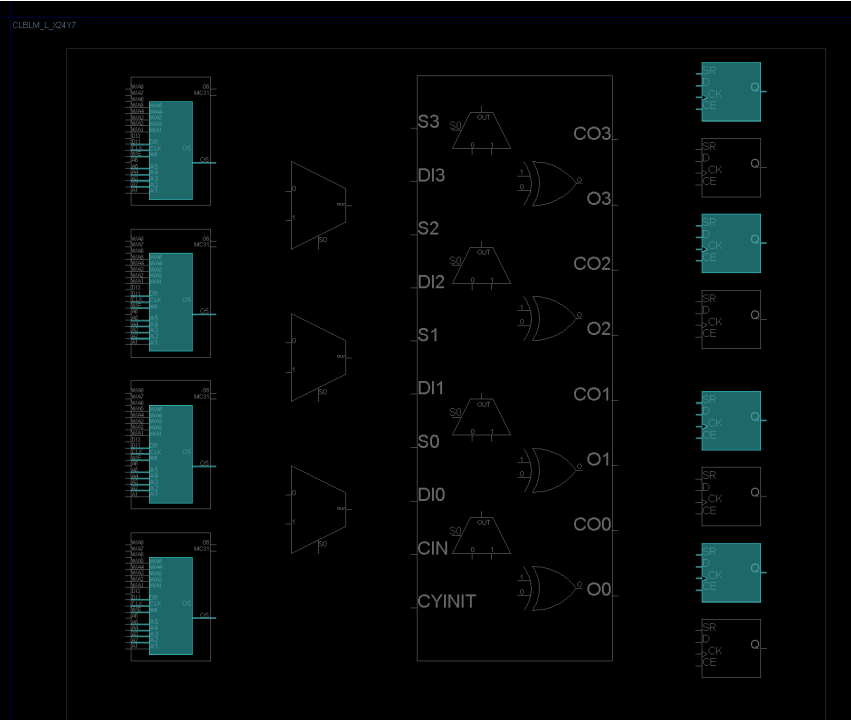

Name	Slack	Levels	High Fanout	From	Total Delay	To	Logic Delay	Net Delay	Requirement	Source Clock	Destination Clock	Exception	Clock Uncertainty
Path 1	2.081	10	1	filter/mul_out[1]CLK	21.412	filter/result_0[PCIN]0	21.336	0.076	25.0	clock	clock		0.035
Path 2	2.081	10	1	filter/mul_out[1]CLK	21.412	filter/result_0[PCIN]10	21.336	0.076	25.0	clock	clock		0.035
Path 3	2.081	10	1	filter/mul_out[1]CLK	21.412	filter/result_0[PCIN]11	21.336	0.076	25.0	clock	clock		0.035
Path 4	2.081	10	1	filter/mul_out[1]CLK	21.412	filter/result_0[PCIN]12	21.336	0.076	25.0	clock	clock		0.035
Path 5	2.081	10	1	filter/mul_out[1]CLK	21.412	filter/result_0[PCIN]13	21.336	0.076	25.0	clock	clock		0.035
Path 6	2.081	10	1	filter/mul_out[1]CLK	21.412	filter/result_0[PCIN]14	21.336	0.076	25.0	clock	clock		0.035
Path 7	2.081	10	1	filter/mul_out[1]CLK	21.412	filter/result_0[PCIN]15	21.336	0.076	25.0	clock	clock		0.035
Path 8	2.081	10	1	filter/mul_out[1]CLK	21.412	filter/result_0[PCIN]16	21.336	0.076	25.0	clock	clock		0.035
Path 9	2.081	10	1	filter/mul_out[1]CLK	21.412	filter/result_0[PCIN]17	21.336	0.076	25.0	clock	clock		0.035
Path 10	2.081	10	1	filter/mul_out[1]CLK	21.412	filter/result_0[PCIN]18	21.336	0.076	25.0	clock	clock		0.035

Figure 9: Critical path

4.3 Area Occupation

In the following, there are the screenshots that show the area occupied by the filter





4.4 Power Summary

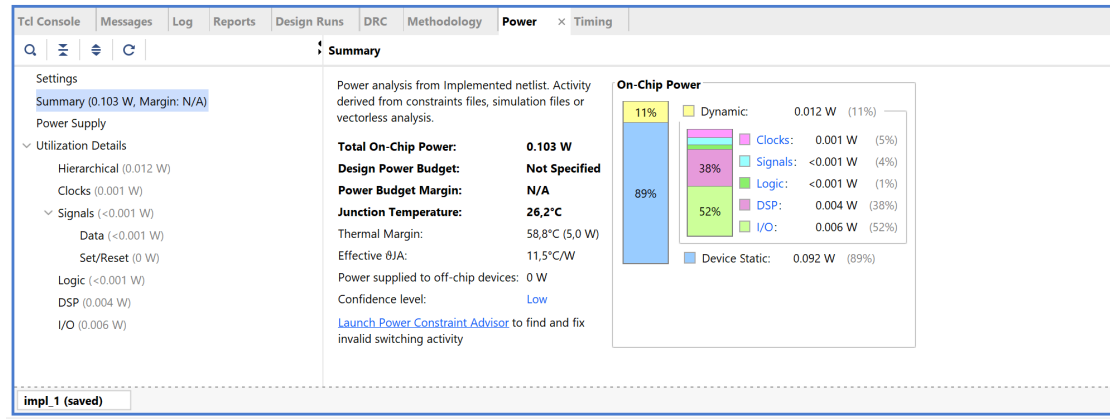


Figure 10: Power Summary

4.5 Utilization Summary

Resource	Utilization	Available	Utilization %
LUT	25	17600	0.14
LUTRAM	16	6000	0.27
FF	216	35200	0.61
DSP	12	80	15.00
IO	34	100	34.00

Figure 11: Utilization Summary

5 Conclusions

In conclusion, I realized a FIR filter with a SRRC impulsive response that takes as an input a signal represented on 16 bits and returns as an output a 16 bits signal. The checks carried out demonstrated that the filter performs well, with a very good precision. The maximum usable frequency could be incremented by dividing the combinatorial network in different sections interleaved by flip-flops, in order to obtain a pipeline with different stages, but this lies outside from the objective of this project and I preferred a lower frequency, in order to have an output that is not too much delayed with the respect of the related input.